


Practice Test Review

Review 29e6e023-72f8-461a-9c1e-84108ac7...

NVIDIA : NCP-AAI

Overall Test Result



The practice test includes only multiple-choice questions (MCQs) due to development limitations. To achieve a good score, you need to review the premium PDF.

All Correct Only Incorrect Only

Question: 1

CertyIQ

When designing tool integration for an agent that needs to perform mathematical calculations, web searches, and API calls, which architecture pattern provides the most scalable and maintainable approach?

☐ External tool services with manual configuration for each agent instance

☒ Microservice-based tool architecture with standardized interfaces ✓

☐ Monolithic tool handler with conditional logic for different tool types

☐ Embedded tool functions within the main agent code

Show Explanation

Question: 2

CertyIQ

A company is deploying an AI-powered customer support agent that integrates external APIs and handles a wide range of customer inputs dynamically.

Which of the following strategies are appropriate when designing an AI agent for dynamic conversation management and external system

interaction? (Choose two.)

- ☒ Integrating a feedback loop from user interactions to iteratively improve agent behavior. ✓
- ☐ Using rule-based logic as the primary framework to maintain consistency in agent decisions.
- ☒ Implementing retry logic for API failures to ensure robustness in external communications. ✓
- ☐ Preferring hardcoded responses for frequent queries to deliver reliable and low-latency answers.

Show Explanation

Question: 3

CertyIQ

In the context of agent development, how does an autonomous agent differ from a predefined workflow when applied to complex enterprise tasks?

- ☐ Agents optimize for execution speed under fixed input-output mappings, while workflows prioritize goal alignment through adaptive reasoning and memory mechanisms.
- ☒ Workflows provide deterministic task sequencing with conditional branching, while agents adapt decisions dynamically based on goals, context, and environment feedback. ✓
- ☐ Workflows emphasize parallelism and distributed coordination of processes, while agents emphasize serialization and isolated problem solving.

Show Explanation

Question: 4

CertyIQ

A Lead AI Architect at a global financial institution is designing a multi-agent fraud detection system using an agentic AI framework. The system must operate in real time, with distinct agents working collaboratively to monitor and analyze transactional patterns across accounts, retain and share contextual information over time, and escalate suspicious behaviors to a human fraud analyst when needed. Which architectural approach enables intelligent specialization, shared memory, and inter-agent coordination in a dynamic and evolving threat environment?

- ☒ Design a modular multi-agent system where individual agents collaborate asynchronously using shared memory and structured messaging. ✓
- ☐ Design a multi-agent system where individual agents collaborate synchronously using shared memory and structured messaging.
- ☐ Design a centralized rule-based service that checks all transactions against static fraud indicators and sends alerts when thresholds are exceeded.
- ☐ Design an agentic workflow where each agent acts independently on isolated data slices with no inter-agent communication to reduce latency and model complexity.
- ☐ Design monolithic LLM-based agents that handle all fraud detection tasks within a single loop, without modular roles or multi-agent coordination.

Show Explanation

Question: 5

CertyIQ

When evaluating coordination failures in a multi-agent system managing distributed manufacturing workflows, which analysis approach best identifies state management and planning synchronization issues?

- ☐ Monitor agent outputs individually to confirm local correctness and examine results of specific workflow steps.
- ☒ Deploy distributed state tracing across agents, analyze transition timing, study communication overhead, and verify synchronization accuracy. ✓
- ☐ Assess synchronization methods during design reviews and use simulations to evaluate coordination across representative workflow scenarios.
- ☐ Track workflow throughput and task completions to measure performance trends and highlight workflow outcomes.

Show Explanation

Question: 6

CertyIQ

You are designing an AI agent for summarizing medical documents that include images and text as well. It must extract key information and recognize dates.
Which feature is most critical for ensuring the agent performs well across multiple input and output formats?

- ☐ Use of guardrails to filter out hallucinated content
- ☐ Retry logic implementation to ensure robustness during API failures
- ☐ Chain-of-thought prompting for reasoning accuracy
- ☒ Multi-modal model integration to handle both text and vision inputs ✓

Show Explanation

Question: 7

CertyIQ

An AI engineer is evaluating an underperforming multi-agent workflow built with NVIDIA agentic frameworks.
Which analysis approach most effectively identifies optimization opportunities in agent coordination and communication patterns?

- ☐ Monitor workflow completion times using analysis that subsumes inter-agent communication costs, coordination overhead, and task allocation balance.
- ☐ Focus exclusively on individual agent accuracy without analyzing workflow-level efficiency, coordination costs, or overall system throughput.
- ☐ Evaluate agents individually, allowing the toolkit to automatically infer interaction effects, communication patterns, and emergent behaviors from coordination.
- ☒ Trace agent interaction patterns using observability features, measure communication overhead, identify redundant operations, and analyze task distribution efficiency. ✓

Show Explanation

Question: 8

CertyIQ

You are designing a virtual assistant that helps users check weather updates via external APIs. During testing, the agent frequently calls the incorrect tools, often hallucinating endpoints or returning incorrect formats. You suspect the prompt structure might be the root cause of these failures.
Which prompt design best supports consistent tool invocation in this agent?

- ☐ Rely on the agent’s internal knowledge to infer tool usage
- ☐ Include tool names in natural language but without parameter examples
- ☐ Provide only a generic system instruction with no examples
- ☒ Use structured prompt templates with few-shot tool usage examples ✓

Show Explanation

Question: 9



You’re working with an LLM to automatically summarize research papers. The summaries often omit critical findings. What’s the best way to ensure that the summaries accurately reflect the core insights of the research papers?

- ☐ Asking the LLM to “summarize the paper.”
- ☐ Asking the LLM to “understand” the paper to generate a summary.
- ☐ Having the LLM generate the summaries and then manually review every output.
- ☒ Asking the LLM to “extract the key findings.” ✓

Show Explanation

Question: 10



A development team is building a customer support agent that interacts with users via chat. The agent must reliably fetch information from external databases, handle occasional API failures without crashing, and improve its responses by learning from user feedback over time. Which of the following tasks is most critical when enhancing an AI agent to handle real-world interactions and improve over time?

- ☐ Applying a well-structured training process with foundational generative models and prompt engineering
- ☐ Utilizing internal knowledge bases to support agent responses alongside external APIs
- ☒ Implementing retry logic for error handling and integrating user feedback loops for iterative improvement ✓
- ☐ Designing conversation flows that provide consistent responses based on predefined scripts

Show Explanation

Question: 11



You are evaluating your RAG pipeline. You notice that the LLM-as-a-Judge consistently assigns high similarity scores to responses that contain irrelevant information. What should you investigate as the most likely potential cause with the least development effort?

- ☐ The temperature setting used by the LLM during response generation.
- ☐ The size of the knowledge base used to power the RAG pipeline.
- ☐ The quality of the synthetic questions used for evaluation.
- ☒ The prompt used to instruct the LLM-as-a-Judge to assess the response. ✓

Show Explanation

Question: 12



A customer service agentic AI is designed to resolve billing inquiries. It consistently resolves inquiries accurately and efficiently. However, a significant number of customers are reporting frustration due to the agent’s tendency to repeatedly ask for the same information (account number, address) during each interaction, even after it’s already been provided. Which evaluation method would be most effective for addressing this issue?

- ☐ Adjusting the agent’s reward function to prioritize speed of resolution over customer satisfaction.
- ☒ Analyzing the agent’s dialogue transcripts to identify patterns in its questioning techniques. ✓
- ☐ Implementing a “conversational flow” analysis to optimize the order of questions asked during each interaction.
- ☐ Increasing the agent’s processing speed to reduce the time it takes to handle each inquiry and increase customer satisfaction.

Show Explanation

Question: 13

CertyIQ

A financial services agentic AI is being used to automate initial customer onboarding. The agent is completing the process efficiently and accurately, but reviews of its conversations reveal it often uses overly formal and complex language that confuses customers. Which type of evaluation is best suited to address this issue?

- ☒ Controlled user testing sessions to collect user feedback on the clarity and tone of responses ✓
- ☐ Compliance review of the agent’s access to regulatory guidelines and policy documentation
- ☐ Continuous user feedback collection, specifically gathering subjective assessments of the agent’s communication style
- ☐ Statistical analysis of the agent’s decision-making patterns to detect overly formal and complex response choices

Show Explanation

Question: 14

CertyIQ

You’re evaluating the performance of a tool-using agent (e.g., one that issues API calls or executes functions). From the list below, what are two important features to evaluate? (Choose two.)

- ☒ Tool use accuracy ✓
- ☐ Tokens per second
- ☐ Tool use rate
- ☒ Task completion rate ✓

Show Explanation

Question: 15

CertyIQ

You’re evaluating the RAG pipeline by comparing its responses to synthetic questions. You’ve collected a large set of similarity scores. What’s the primary benefit of aggregating these scores into a single metric (e.g., average similarity)?

- ☐ Aggregation identifies the specific chunks within the RAG pipeline that are contributing to the highest similarity scores.
- ☒ Aggregation reduces the complexity of the evaluation process and allows for a more overall assessment of the pipeline’s effectiveness. ✓
- ☐ Aggregation provides a more accurate representation of the RAG pipeline’s performance.
- ☐ Aggregation eliminates the need for qualitative analysis of the RAG pipeline’s responses.

Show Explanation

Question: 16

CertyIQ

In designing an AI workflow which of the following best describes a comprehensive approach to improving the performance of AI agents?

- ☐ Implementing benchmarking pipelines, deploying physical agents and monitoring user engagement metrics
- ☒ Implementing benchmarking pipelines, collecting user feedback, and tuning model parameters iteratively ✓
- ☐ Implementing benchmarking pipelines and incorporating a dynamic dataset for a real-time fall-back
- ☐ Monitoring agents’ throughput and time-to-first-token from the scoring engine

Show Explanation

Question: 17

CertyIQ

A team is evaluating multiple versions of an AI agent designed for customer support. They want to identify which version completes tasks more efficiently, responds accurately, and improves over time using user feedback. Which practice is most important to ensure continuous refinement and optimal performance of the AI agent?

- ☐ Comparing agents on isolated tasks without standardized benchmarking pipelines
- ☐ Relying solely on offline benchmarks without incorporating live user feedback during tuning
- ☒ Implementing an evaluation framework that quantifies task efficiency and incorporates human-in-the-loop feedback ✓
- ☐ Tuning model parameters once before deployment to maximize initial accuracy

Show Explanation

Question: 18

CertyIQ

When analyzing inconsistent performance across a fleet of customer service agents handling similar queries, which evaluation approach most effectively identifies root causes and optimization opportunities?

- ☐ Assess performance data from recently improved agents and highlight strong results, using outcome comparisons to identify areas with the greatest impact on service quality.
- ☐ Average performance metrics across all agents as this will smooth individual variations, query distribution differences, and temporal factors affecting agent behavior and accuracy.
- ☒ Deploy stratified evaluation sampling across agent variants, query complexity levels, and temporal patterns while tracking decision paths using comparative analytics. ✓
- ☐ Review performance across both high- and low-accuracy agent groups, comparing case outcomes and identifying patterns contributing to top and bottom results.

Show Explanation

Question: 19

CertyIQ

Your agent is generating inconsistent and contradictory statements. Which approach would be most suitable to improve the agent’s output?

- ☒ Employing Reflexion ✓
- ☐ Increasing the number of generated plans
- ☐ Using Decomposition-First Planning
- ☐ Decreasing the length of prompts

Show Explanation

Question: 20

CertyIQ

You’re utilizing an LLM to translate complex technical documentation into multiple languages. The translations often lack nuance and fail to capture the original intent. What’s the most effective strategy for improving the quality of the translations?

- ☒ Providing the LLM with a glossary of key terms, concepts in all languages and the dataset of previously translated text. ✓
- ☐ Training the LLM on a dataset of translated texts.
- ☐ Providing the LLM with guidance to “translate the documents” without additional guidance, so it can use trained knowledge.
- ☐ Providing the LLM with guidance to translate “with high accuracy” without additional guidance, so it can use trained knowledge.

Show Explanation

Question: 21



An e-commerce platform is implementing an AI-powered customer support system that handles inquiries ranging from simple FAQ responses to complex product recommendations and technical troubleshooting. The system experiences unpredictable traffic patterns with sudden spikes during sales events and varying complexity requirements. Simple questions comprise the majority of requests but require minimal compute, while complex product recommendations need sophisticated reasoning. The company wants to optimize costs while maintaining service quality across all query types. Which approach would provide the MOST cost-optimized scaling strategy for this variable-workload, mixed-complexity environment?

- ☐ Deploy specialized NVIDIA NIM microservices using a single large model configuration that handles all agent functions on high-capacity GPUs, with auto-scaling infrastructure that maintains constant resource allocation across all traffic patterns.
- ☐ Deploy specialized NVIDIA NIM microservices on CPU-optimized infrastructure with auto-scaling capabilities to minimize hardware costs, while accepting longer inference times for cost optimization benefits.
- ☒ Deploy specialized NVIDIA NIM microservices with an LLM router to dynamically route requests to appropriate models based on complexity, combined with auto-scaling infrastructure that scales different model types independently. ✓
- ☐ Deploy multiple specialized NVIDIA NIM microservices with identical high-capacity models across all available GPUs, implementing auto-scaling infrastructure without request complexity differentiation or dynamic model selection capabilities.

Show Explanation

Question: 22



A technology startup is preparing to launch an AI agent platform to serve clients with unpredictable usage patterns. They face periods of high user activity and low demand, so their deployment approach must minimize wasted resources during slow times and automatically allocate more resources during busy periods – all while keeping operational costs reasonable. Given these requirements, which deployment strategy most effectively ensures both cost-effectiveness and adaptability for scaling agentic AI systems?

- ☐ Scheduling periodic manual reviews to increase or decrease infrastructure based on predicted user numbers
- ☐ Monitoring system logs for usage patterns and making infrastructure changes after monthly analysis
- ☐ Using fixed-size virtual machine clusters to guarantee consistent resource allocation at all times
- ☒ Implementing autoscaling policies in a container orchestration environment to automatically adjust resources according to workload changes ✓

Show Explanation

Question: 23



A company operates agent-based workloads in multiple data centers. They want to minimize latency for users in different regions, maintain continuous service during infrastructure upgrades, and keep operational costs predictable. Which deployment practice best supports low-latency, resilient, and cost-efficient agent operations at scale?

- ☐ Schedule regular agent downtime for system updates and operational recalibration.
- ☒ Implement geo-distributed deployments with rolling updates and resource usage monitoring. ✓
- ☐ Prioritize high-performance GPUs for all agents in geo-distributed deployments.
- ☐ Apply static infrastructure allocation with centralized resource usage monitoring at a single data center.

Show Explanation

Question: 24

CertyIQ

In a global financial firm, an AI Architect is building a multi-agent compliance assistant using an agentic AI framework. The system must manage short-term memory for multi-turn interactions and long-term memory for persistent user and policy context. It should enable contextual recall and adaptation across sessions using NVIDIA's tool stack. Which architectural approach best supports these requirements?

- ☒ Leverage NVIDIA NeMo Framework with modular memory management, integrating conversational state tracking, knowledge graphs, and vector store retrieval, while using LoRA-tuned models to adapt responses overtime. ✓
- ☐ Leverage RAPIDS cuDF for memory tracking by streaming multi-turn conversation logs as GPU-resident data frames, assuming transactional history can be recalled and reasoned over using dataframe operations.
- ☐ Rely exclusively on TensorRT to encode all prior knowledge into compiled model weights, allowing inference-only execution with no external memory dependencies across sessions.
- ☐ Leverage NVIDIA Triton Inference Server with dynamic batching to cache session-level inputs between inference calls, and use an external Redis store for long-term memory.

Show Explanation

Question: 25

CertyIQ

A development team is building an AI agent capable of autonomously planning and executing multi-step tasks while retaining context and learning from past interactions. Which practice is most important to enable the agent to effectively manage long-term memory and complex tasks?

- ☒ Implement memory mechanisms for context retention and apply chain-of-thought prompts to enhance reasoning. ✓
- ☐ Use basic rule-based decision methods that emphasize fast responses over adaptive planning.
- ☐ Apply short-term memory approaches that handle each interaction independently of previous ones.
- ☐ Reduce planning features and memory management to keep the system streamlined.

Show Explanation

Question: 26

CertyIQ

You are developing an agent that needs to perform a complex set of tasks repeatedly. Why is periodic fine-tuning an important aspect of long-term knowledge retention for this type of agent?

- ☐ It prevents the agent from becoming overly specialized to a single task.
- ☐ It eliminates the need for external storage like RAG.
- ☒ It prevents the agent from forgetting past successes and failures. ✓
- ☐ It guarantees the agent will produce the same output for the same input.

Show Explanation

Question: 27

CertyIQ

An agent is tasked with solving a series of complex mathematical problems that require external tools to find information. It often struggles to keep track of intermediate steps and reasoning. Which prompting technique would be MOST effective in improving the agent's clarity and reducing errors in its reasoning?

- ☒ ReAct ✓
- ☐ Symbolic Planning
- ☐ Zero-shot CoT
- ☐ Multi-Plan Generation

Show Explanation

Question: 28

CertyIQ

When analyzing memory-related performance degradation in agents handling extended customer support sessions, which evaluation methods effectively identify optimization opportunities for context retention? (Choose two.)

- ☐ Clear memory after each interaction and reset session state, removing historical context needed for personalized tasks to identify optimization opportunities.
- ☒ Profile memory access patterns by measuring retrieval latency, relevance scoring accuracy, and storage efficiency while monitoring context window utilization to identify optimization opportunities. ✓
- ☐ Use fixed memory allocation including all conversation types, topic changes, and user needs, allowing adaptive-free observation of interaction patterns to identify optimization opportunities.
- ☒ Implement sliding window analysis comparing context compression strategies, summarization quality, and information preservation rates across varying conversation lengths to identify optimization opportunities. ✓
- ☐ Store all conversation history including all interactions, allowing adaptive-free observation of data to identify optimization opportunities.

Show Explanation

Question: 29

CertyIQ

An enterprise wants their AI agent to support complex project management tasks. The agent should remember ongoing project details, adjust its plans based on new information, and break down large goals into actionable steps. Which strategy best enables the AI agent to autonomously decompose tasks and adapt to new Information over time?

- ☐ Predefining static workflows for each project type to guarantee consistent execution
- ☒ Developing long-term knowledge retention strategies and dynamic state management for adaptive planning ✓
- ☐ Storing recent user interactions in a temporary cache for immediate retrieval
- ☐ Applying rule-based logic to each new request isolated from previous project data

Show Explanation

Question: 30

CertyIQ

In a ReAct (Reasoning-Acting) agent architecture, what is the correct sequence of operations when the agent encounters a complex multi-step problem requiring external tool usage?

- ☐ Thought --> Answer --> Action --> Observation
- ☐ Action --> Thought --> Observation --> Action --> Thought --> Observation --> Answer
- ☐ Observation --> Thought --> Action --> Observation --> Thought --> Action --> Answer
- ☒ Thought --> Action --> Observation --> Thought --> Action --> Observation --> Answer ✓

Show Explanation

Question: 31



What is a key limitation of Chain-of-Thought (CoT) prompting when using smaller language models for reasoning tasks?

- ☐ CoT prompting simplifies error analysis for small models, making it easy to identify and correct mistakes at each reasoning step.
- ☐ CoT prompting ensures step-by-step outputs, enabling even small models to solve complex problems reliably.
- ☒ CoT prompting requires relatively large models; smaller models may produce reasoning chains that appear logical but are actually incorrect, leading to poorer performance. ✓
- ☐ CoT prompting consistently improves the logical accuracy of outputs for both small and large language models.

Show Explanation

Question: 32



In a production agentic system handling thousands of concurrent conversations, which state management strategy provides optimal performance while ensuring context preservation?

- ☐ Global shared state with locks for concurrent access
- ☒ Session-isolated state with serialization and lazy loading ✓
- ☐ Stateless design with context reconstruction from message history

Show Explanation

Question: 33



Your agent’s primary task is to collect data from a device management platform API. This API frequently returns data in JSON format, but the underlying schemas can evolve without detailed documentation. What’s the most effective strategy for the agent to maintain consistent access to the data, despite potential schema changes?

- ☐ Building a JSON schema validator that constantly monitors and enforces the current API schema.
- ☐ Manually inspecting the API response’s structure and updating the agent’s code when a change is detected.
- ☐ Building a fixed data extraction method based on the API’s historical response patterns.
- ☒ Building a flexible data mapping layer that can adapt to changes in the API response structure. ✓

Show Explanation

Question: 34



You are improving your Retrieval-Augmented Generation (RAG) pipeline, and you want to exploit the properties of the embedding model to improve the solutions. Within a RAG system, why is reranking retrieved chunks a crucial step after initial retrieval?

- ☐ Reranking improves retrieval accuracy, regardless of the initial retrieval method.
- ☐ Reranking’s purpose is to reduce the overall length of the response generated by the LLM.
- ☒ Reranking allows for prioritizing chunks based on relevance to the semantic intent of the query. ✓
- ☐ Reranking primarily focuses on optimizing the visual presentation of the retrieved chunks.

Show Explanation

Question: 35



Your agent is designed to manage tasks through a service management API. The API responds with detailed event logs, but these logs contain both metadata and structured data. To ensure the agent correctly interprets and processes the data from these logs, what's the most prudent approach?

- ☒ Employ a specialized parser that adheres to the API's documentation, to insure strict adherence to structured data. X
- ☐ Employing a modular design that allows the agent to dynamically adjust its parsing logic. ✓
- ☐ Using a human-in-the-loop approach, manually inspecting and interpreting each log entry.
- ☐ Employ a specialized parser that extracts all data fields, regardless of their type.

Show Explanation

Question: 36



Your deployed legal assistant shows great performance but occasionally repeats incorrect legal terms. Which tuning method best improves factual reliability?

- ☐ Replace retrieval with static hard-coded text snippets
- ☒ Use more verbose prompts to reinforce correct definitions X
- ☐ Increase output randomness to improve exploration
- ☐ Add fact-checking steps using external tools during generation ✓

Show Explanation

Question: 37



An AI Engineer is experimenting with data retrieval performance within a RAG system. Which of the following techniques is most likely to improve the quality of the retrieved chunks?

- ☒ Adding clarifying keywords and synonyms to the original query to broaden the search. ✓
- ☐ Truncating long queries to fit within the LLM's context window.
- ☐ Using a single, highly specific keyword to guarantee a precise match.
- ☐ Directly feeding the original query to the LLM without any modification.

Show Explanation

Question: 38



You're building a RAG system that uses RAG Fusion. Which of the following approaches would be most effective in determining how to combine information from multiple retrieved chunks?

- ☐ Filtering out chunks considered inconsistent with others before presenting information to the LLM.
- ☐ Using the LLM to automatically identify the most important sentences within each chunk and combine them. ✓
- ☐ Manually selecting the most relevant sentences from each chunk and inserting them into the LLM prompt.
- ☒ Concatenating the text from all retrieved chunks into a single block to form the response. X

Show Explanation

Question: 39



You are tasked with deploying a multi-modal agentic system that must respond to user queries with minimal latency while maintaining guardrails for safe and context-aware interactions. Which of the following configurations best leverages NVIDIA's AI stack to meet these requirements?

- ☒ Integrate NeMo Guardrails, configure NIM microservices for optimized inference, use TensorRT-LLM for deployment, and profile the system using Triton Inference Server with multi-modal support. ✓
- ☐ Integrate NeMo Guardrails, use Omniverse to generate synthetic data, configure NIM microservices for optimized inference, use TensorRT-LLM for deployment, and profile the system using NeMo Agent Toolkit for multi-modal support.
- ☐ Use NeMo Guardrails for safety, deploy the model with Triton Inference Server using default settings, and rely on hardware accelerators like GPU/TPU inference for cost efficiency.
- ☐ Use NIM microservices for deployment, optionally use NeMo Guardrails unless one wants to minimize the inference overhead.

Show Explanation

Question: 40



When evaluating optimization opportunities between NeMo Guardrails, NIM microservices, and TensorRT-LLM in a production healthcare agent, which analysis approach best identifies optimization opportunities across the NVIDIA stack?

- ☐ Conduct stress testing of individual microservices and guardrails to measure peak throughput and determine theoretical performance limits of each module.
- ☐ Use default configurations to establish a deployment baseline, focusing on stability before conducting deeper performance profiling.
- ☒ Create end-to-end latency waterfalls that capture guardrail overhead, NIM queuing delays, and TensorRT optimization benefits while assessing overall pipeline efficiency. ✓
- ☐ Tune each component individually, focusing primarily on local performance metrics with secondary attention to integration patterns.

Show Explanation

Question: 41



When evaluating GPU utilization inefficiencies in deploying Llama Nemotron models across A100 and H100 clusters, which approaches help identify optimal resource allocation strategies? (Choose two.)

- ☒ Allow Nemotron variants to profile actual workload characteristics and allocate resources based on observed demands. ✗
- ☐ Profile resource utilization for each Nemotron variant and match models to appropriate GPU tiers. ✓
- ☐ Allocate all agents to H100 GPUs, allowing resource profiles to automatically adjust for model size and computational requirements.
- ☒ Assess concurrent execution capabilities by employing multi-instance GPU partitioning for varying workload types. ✓

Show Explanation

Question: 42



Your team notices a spike in failed tool calls from a deployed workflow agent after a recent API schema update. The agent still returns outputs, but many are irrelevant or incomplete. Which maintenance task should be prioritized to restore accurate behavior?

- ☐ Reset the agent's long-term memory and reinitialize logs.
- ☒ Update the tool function specifications and re-test action sequences. ✓
- ☐ Increase model temperature to encourage tool exploration.

- ☐
- Reduce tool retrieval vector similarity threshold to broaden context.

Show Explanation

Question: 43

CertyIQ

An AI architect at a national healthcare provider is maintaining an agentic AI system. The system must monitor model and system performance in real time, raise alerts on failures or anomalies, manage version control and rollback of diagnostic models, and provide transparent insight into agent behavior during patient care workflows. Which operational approach best supports these requirements using the NVIDIA AI stack?

- ☐
- Containerize each agent in NIM with basic health checks running on cron jobs, and manage version rollback by swapping prebuilt container images.

☐

 Optimize all models with TensorRT and use periodic manual log reviews and NVIDIA shell scripts for detecting service anomalies and managing rollback.

☒

 Deploy agent models on NVIDIA Triton Inference Server with Prometheus and Grafana for performance alerting, and manage model lifecycle via NGC and the Triton model repository. ✓

☐

 Expose agents as stateless NVIDIA API endpoints and monitor activity through application logs, with model versions tracked in a Git-based script repository.

Show Explanation

Question: 44

CertyIQ

Your support agent frequently fails to complete tasks when third-party tools return unexpected formats. Which solution improves resilience against these failures?

- ☒

 Add robust schema validation and exception handling for all tool outputs ✓
- ☐

 Use deterministic temperature settings for all generations
- ☐

 Reduce the number of tools available to avoid bad integrations
- ☐

 Re-train the model to avoid the use of third-party tools entirely

Show Explanation

Question: 45

CertyIQ

A customer service agent sometimes fails to complete multi-step workflows when APIs respond slowly or inconsistently. Which approach most effectively increases robustness when working with unreliable APIs?

- ☐

 Restrict available tools to reduce decision complexity
- ☒

 Add retries with exponential backoff and set request timeouts ✓
- ☐

 Cache recent API results to limit unnecessary repeated calls
- ☐

 Adjust generation parameters to produce more predictable responses

Show Explanation

Question: 46

CertyIQ

You are building an agent that performs financial analysis by retrieving and processing structured data from a client’s internal SQL database. The agent must handle occasional connection errors and retry the query up to a few times before failing gracefully. Which approach best meets these requirements?

- ☒ Use structured tool calls with built-in retry handling and timed delays inside the tool wrapper ✓
- ☐ Use few-shot prompting to guide the agent’s conversation flow and manually retry failed API responses
- ☐ Use a reactive agent pattern that retries the query after a user confirms a retry attempt
- ☐ Use memory to track the number of failed attempts and apply it in later retries

Show Explanation

Question: 47

CertyIQ

After deploying a financial assistant agent, users report occasional inconsistencies in how transactions are categorized. What is the best first step for diagnosing the issue?

- ☐ Review and modify prompt temperature to enhance precision
- ☒ Review and retrain the model with more financial datasets X
- ☐ Implement agent memory reset after each session
- ☐ Review tool call inputs and outputs in recent session logs ✓

Show Explanation

Question: 48

CertyIQ

You are implementing a RAG (Retrieval-Augmented Generation) solution. What is the primary purpose of implementing semantic guardrails within a RAG system?

- ☐ To establish rules and constraints based on the meaning of user queries and generated responses. ✓
- ☐ To eliminate all potential harmful entries from the vector database.
- ☐ To automatically translate all LLM responses into multiple languages for improved user comprehension.
- ☒ To filter out all queries containing specific keywords that have been flagged as problematic. X

Show Explanation

Question: 49

CertyIQ

This question addresses important concerns in the field of AI ethics and compliance, particularly as organizations develop more autonomous AI agents. Implementing effective guardrails against bias, ensuring data privacy, and adhering to regulations are essential components of responsible AI development. Which of the following statements accurately describes how RAGAS (Retrieval Augmented Generation Assessment) can be utilized for implementing safety checks and guardrails in agentic AI applications?

- ☐ RAGAS cannot evaluate all safety aspects independently but provides metrics like Topic Adherence and Agent Goal Accuracy that serve as guardrails. ✓
- ☒ RAGAS can only evaluate the quality of document retrieval but has no applications for safety guardrails in agentic systems. X
- ☐ RAGAS is exclusively designed for hallucination detection and cannot evaluate other safety aspects of agentic applications.
- ☐ RAGAS can only be used in conjunction with other guardrail frameworks like NeMo and cannot function independently.

Show Explanation

Question: 50

CertyIQ

Your team has deployed a generative agent for internal HR use, including summarizing candidate resumes and suggesting interview questions. After deployment, you’ve noticed that the model occasionally associates certain names or genders with particular roles. Which mitigation strategy is the most effective and scalable for reducing this type of bias in agent outputs?

- ☐ Adjust system prompts to explicitly instruct the agent to avoid assumptions based on demographic features
- ☐ Randomly replace names in prompts to reduce identity correlation
- ☐ Add more training examples to the training dataset and re-train the model
- ☒ Implement guardrails to prevent outputs referencing protected attributes ✓

Show Explanation

Question: 51

CertyIQ

When analyzing safety violations in a financial advisory agent that uses NeMo Guardrails, which evaluation approach best identifies gaps in guardrail coverage?

- ☐ Apply keyword-and rule-based validation methods to confirm compliance with policy terms and common risk conditions.
- ☒ Analyze violation patterns, test adversarial prompts, measure guardrail activation, and align policies with observed failures. ✓
- ☐ Conduct functional testing with representative user inputs to verify policy enforcement in typical usage scenarios.
- ☐ Monitor overall guardrail activations and system logs to assess operational behavior across different interaction types.

Show Explanation

Question: 52

CertyIQ

You’re deploying a healthcare-focused agentic AI system that helps doctors make treatment recommendations based on patient records. The agent’s reasoning is not exposed to users, and its decisions sometimes differ from clinical guidelines. What safety and compliance mechanisms should be in place? (Choose two.)

- ☒ Allow overrides by human doctors to maintain accountability ✓
- ☒ Require model explainability or traceability for all outputs ✓
- ☐ Prioritize autonomous speed of decision over explainability
- ☐ Exempt the model from compliance if it improves outcomes
- ☐ Obfuscate decision logic to protect proprietary methods

Show Explanation

Question: 53

CertyIQ

An AI Engineer is analyzing a production agentic AI system’s compliance with responsible AI standards. Which evaluation approaches effectively identify potential safety vulnerabilities and ethical risks in multi-agent workflows? (Choose two.)

- ☐ Emphasize latency metrics and throughput performance as key evaluation factors for safety vulnerabilities, providing a baseline for operational measures and resource allocation.
- ☒ Implement comprehensive audit trails using NVIDIA NeMo Guardrails with semantic similarity checks, tracking agent decisions across conversation flows and evaluating policy violations through automated compliance scoring. ✓
- ☐ Use user feedback as a primary signal for risk identification, emphasizing post-deployment observations and qualitative experience reports alongside operational monitoring.
- ☒ Deploy multi-layered evaluation combining bias detection metrics (demographic parity, equalized odds) with adversarial testing to probe agent responses for harmful outputs across diverse user populations ✓

Show Explanation

Question: 54

CertyIQ

You are deploying an AI-driven applicant-screening agent that analyzes candidate resumes and social-media data to recommend top applicants. Due to anti-discrimination laws and corporate policy, the system must mitigate bias against protected groups, maintain an audit trail of decisions, and comply with GDPR (including data minimization and explicit consent). Which of the following strategies is most effective for ensuring your screening agent both mitigates bias in its recommendations and complies with data-privacy regulations?

- ☐ Perform a post-deployment GDPR and bias audit and process raw personal data as received.
- ☒ Pseudonymize protected attributes, implement fairness-aware debiasing, maintain an audit trail, and enforce GDPR data-minimization and consent. ✓
- ☐ Encrypt all candidate data at rest and in transit, remove protected attributes from analysis, and conduct manual bias checks on recommendations.
- ☐ Exclude gender and ethnicity fields during training, use a generic privacy policy for consent, and do not maintain audit logs or apply targeted debiasing.

Show Explanation

Question: 55

CertyIQ

A health assistant agent has been running on production environment for several weeks. The compliance team wants to audit how personal health data has been processed. Which operational feature supports this requirement?

- ☐ Adding more prompt examples to clarify privacy rules
- ☐ Masking all output with a profanity and PII detector
- ☐ Increasing model temperature for diverse interpretations
- ☒ Enabling full session logging with audit trail metadata ✓

Show Explanation

Question: 56

CertyIQ

An AI Engineer has deployed a multi-agent system to manage supply chain logistics. Stakeholders request greater insight into how the agents decide on actions across tasks. Which approach would best improve decision transparency without modifying the underlying model architecture?

- ☐ Gather structured user evaluations after each completed subtask
- ☐ Generate visual summaries of attention patterns for every decision
- ☒ Record a step-by-step reasoning log throughout each agent workflow ✓
- ☐ Retain and share the full sequence of task instructions with stakeholders

Show Explanation

Question: 57

CertyIQ

When analyzing a customer service agentic system’s performance degradation over time, which evaluation approach most effectively identifies opportunities for human-in-the-loop intervention to improve agent decision-making transparency and user trust?

- ☐ Monitor only final task completion rates without examining intermediate decision points, user interaction patterns, or opportunities for beneficial human intervention during agent conversations
- ☒ Implement multi-stage evaluation tracking decision confidence scores, user correction patterns, intervention effectiveness, and explainability-satisfaction correlations ✓
- ☐ Rely on periodic manual reviews of random conversation samples without systematic tracking of intervention effectiveness, decision transparency, or user trust indicators
- ☐ Collect anonymous usage statistics without capturing specific decision rationales, user feedback on agent explanations, or transparency improvement opportunities for trust building

Show Explanation

Question: 58

CertyIQ

A medical diagnostics company is deploying an agentic AI system to assist radiologists in analyzing medical imaging. The system must provide AI-generated preliminary diagnoses and allow radiologists to review, modify, and approve all recommendations before patient treatment decisions. Human expertise should remain central, with detailed records of human interventions and decision rationales maintained. Which approach would best balance human oversight with AI support in a safety-critical setting?

- ☒ Design an interactive system that presents AI analysis with confidence scores, allows radiologists to review evidence, modify recommendations, and requires explicit approval with documented reasoning for all decisions. ✓
- ☐ Design a fully automated system that presents final diagnoses to radiologists for simple approval or rejection, minimizing human interaction to improve efficiency and reduce decision fatigue.
- ☐ Design a passive monitoring system where AI makes decisions while humans observe without ability to intervene, focusing on post-decision evaluation and quality assurance.
- ☐ Design a simple notification system that alerts radiologists only when AI confidence falls below predetermined thresholds, otherwise allowing autonomous operation without human review or documentation.

Show Explanation

Question: 59

CertyIQ

A development team is building an AI agent capable of autonomously planning and executing multi-step tasks while retaining context and learning from past interactions. Which practice is most important to enable the agent to effectively manage long-term memory and complex tasks?

- ☐ Use basic rule-based decision methods that emphasize fast responses over adaptive planning.
- ☐ Apply short-term memory approaches that handle each interaction independently of previous ones.
- ☒ Implement memory mechanisms for context retention and apply chain-of-thought prompts to enhance reasoning. ✓
- ☐ Reduce planning features and memory management to keep the system streamlined.

Show Explanation

Question: 60

CertyIQ

You are implementing a RAG (Retrieval-Augmented Generation) solution. What is the primary purpose of implementing semantic guardrails within a RAG system?

- ☒ To filter out all queries containing specific keywords that have been flagged as problematic. ✗
- ☐ To automatically translate all LLM responses into multiple languages for improved user comprehension.
- ☐ To eliminate all potential harmful entries from the vector database.
- ☐ To establish rules and constraints based on the meaning of user queries and generated responses. ✓

Show Explanation

Question: 61

CertyIQ

A customer service agentic AI is designed to resolve billing inquiries. It consistently resolves inquiries accurately and efficiently. However, a significant number of customers are reporting frustration due to the agent's tendency to repeatedly ask for the same information (account number, address) during each interaction, even after it's already been provided. Which evaluation method would be most effective for addressing this issue?

- ☐ Adjusting the agent's reward function to prioritize speed of resolution over customer satisfaction.
- ☐ Implementing a "conversational flow" analysis to optimize the order of questions asked during each interaction.
- ☒ Analyzing the agent's dialogue transcripts to identify patterns in its questioning techniques. ✓
- ☐ Increasing the agent's processing speed to reduce the time it takes to handle each inquiry and increase customer satisfaction.

Show Explanation

Question: 62

CertyIQ

An agent is tasked with solving a series of complex mathematical problems that require external tools to find information. It often struggles to keep track of intermediate steps and reasoning. Which prompting technique would be MOST effective in improving the agent's clarity and reducing errors in its reasoning?

- ☒ ReAct ✓
- ☐ Symbolic Planning
- ☐ Zero-shot CoT
- ☐ Multi-Plan Generation

Show Explanation

Question: 63

CertyIQ

A logistics company is implementing an agentic AI system for supply chain optimization that manages inventory levels, predicts demand, and automatically reorders supplies across multiple warehouses. Supply chain managers need to monitor AI decisions, understand the reasoning behind inventory recommendations, and intervene when business conditions change rapidly. The system must present complex data analytics in an intuitive way that enables quick decision-making while providing detailed insights when needed. Managers have varying levels of technical expertise and need interfaces that support both high-level oversight and detailed analysis. Which user interface design approach would BEST support effective human oversight of this complex multi-agent supply chain system?

- ☒ Develop a comprehensive dashboard with AI decision summaries, drill-down access to underlying data sets, and segmented performance metrics to enable targeted analysis of supply chain operations. ✗
- ☐ Create a layered interface featuring intuitive summaries, drill-down capabilities for detailed analysis, contextual explanations of AI decisions, and clear intervention controls with impact visualization and decision support tools. ✓
- ☐ Create a streamlined interface presenting only high-level AI decisions and simplified recommendations, with drill-down views limited to basic historical trends for quick reference.

Show Explanation

Question: 64

CertyIQ

In designing an AI workflow which of the following best describes a comprehensive approach to improving the performance of AI agents?

- ☐ Monitoring agents' throughput and time-to-first-token from the scoring engine
- ☐ Implementing benchmarking pipelines, deploying physical agents and monitoring user engagement metrics

- ☐ Implementing benchmarking pipelines and incorporating a dynamic dataset for a real-time fall-back
- ☒ Implementing benchmarking pipelines, collecting user feedback, and tuning model parameters iteratively ✓

Show Explanation

Question: 65

CertyIQ

When evaluating an agent's integration with external tools and APIs for data retrieval and action execution, which analysis approaches effectively identify reliability and performance issues?
Pick the 2 correct responses below

- ☐ Connect to external APIs with standard procedures and monitor request and response exchanges to isolate the analysis of integration reliability and effectiveness. ✓
- ☒ Design integration tests simulating API version changes, schema modifications, and backward compatibility scenarios to ensure reliable tool connections across updates. ✓
- ☐ Use static API endpoints and parameters configured during development, allowing consistent and effective agent integration across predictable workflows.
- ☒ Implement comprehensive API call tracing with latency measurement, success rates per endpoint, and correlation analysis between tool failures and task completion. ✗

Show Explanation

Question: 66

CertyIQ

A company is building an AI agent that must retrieve information from large document collections and client databases in real time. The team wants to ensure fast, accurate retrieval and maintain high data quality.
Which approach best supports efficient knowledge integration and effective data handling for such an agent?

- ☐ Integrating client data sources as they already incorporate data quality checks or augmentation to speed up deployment
- ☐ Using traditional relational databases because they don't need specialized retrieval mechanisms for all data queries
- ☐ Relying on pre-trained models instead of connecting to external knowledge sources during inference
- ☒ Implementing retrieval-augmented generation (RAG) pipelines combined with vector databases to accelerate access to relevant information ✓

Show Explanation

Question: 67

CertyIQ

Your team notices a spike in failed tool calls from a deployed workflow agent after a recent API schema update. The agent still returns outputs, but many are irrelevant or incomplete.
Which maintenance task should be prioritized to restore accurate behavior?

- ☐ Reset the agent's long-term memory and reinitialize logs.
- ☒ Update the tool function specifications and re-test action sequences. ✓
- ☐ Increase model temperature to encourage tool exploration.
- ☐ Reduce tool retrieval vector similarity threshold to broaden context.

Show Explanation

Question: 68

CertyIQ

An AI Engineer at a retail company is developing a customer support AI agent that needs to handle multi-turn conversations while keeping track of customers' previous queries, preferences, and unresolved issues across multiple sessions.
Which approach is most effective for managing context retention and enabling the agent to respond coherently in real time?

- ☐ Use a sliding window of recent conversation tokens in memory to track only the last few exchanges.
- ☐ Retrain the model periodically using historical logs to improve long-term contextual understanding.
- ☒ Implement a hybrid memory system with vector-based search and key-value storage to retrieve relevant past interactions. ✓
- ☐ Increase the maximum context window size so the full conversation history is processed each time.

Show Explanation

Question: 69

CertyIQ

An autonomous vehicle company operates a multi-agent AI system across its fleet to process real-time sensor data, make driving decisions, and communicate with cloud infrastructure. The company needs fleet-wide monitoring to track GPU utilization, inference times, and memory usage, correlate performance with driving conditions and system load, and predict safety issues before they occur. Which monitoring and observability approach would BEST meet these fleet-scale, safety-critical requirements?

- ☐ Implement layered application monitoring with distributed tracing, synthetic transaction monitoring, and custom dashboards to capture complex dependencies, transaction flow, and service-level performance trends across the fleet.
- ☒ Deploy NVIDIA NIM microservices with Prometheus integration, NVIDIA Nsight Systems profiling, and Kubernetes-native monitoring to provide detailed metrics, profiling, and container orchestration observability across the entire stack. ✗
- ☐ Deploy enterprise telemetry using OpenTelemetry standards with machine learning-based anomaly detection, custom performance visualization, and automated alerting to deliver predictive operational insights and support proactive maintenance actions. ✓
- ☐ Implement comprehensive APM solutions with real-time baselines, automated root cause analysis, and fleet management integration to coordinate operational insights and performance management across thousands of vehicles.

Show Explanation

Question: 70

CertyIQ

When evaluating a multi-agent customer service system experiencing unpredictable scaling costs and performance bottlenecks during peak hours, which analysis approaches effectively identify optimization opportunities for both infrastructure efficiency and service reliability? Pick the 2 correct responses below

- ☐ Deploy agents with configurable scaling workflows, allowing analysis of resource adjustment strategies and their effects on service stability during variable demand periods. ✓
- ☐ Scale agent infrastructure based on aggregate performance trends, using system-wide monitoring tools to identify broader optimization patterns across resources.
- ☒ Implement comprehensive workload profiling using NVIDIA Nsight to analyze GPU utilization patterns, identify underutilized resources, and optimize batch sizing for dynamic scaling with Kubernetes HPA. ✓
- ☒ Deploy distributed tracing with cost attribution per agent type, correlating resource consumption with business value metrics to identify optimization opportunities in agent deployment strategies. ✗
- ☐ Maintain consistent resource allocation across all service hours, for a more precise view of baseline traffic impact on long-term infrastructure efficiency.

Show Explanation

Question: 71

CertyIQ

Which two error handling strategies are MOST important for maintaining agent reliability in production environments? Pick the 2 correct responses below

- ☒ Automatic retry with exponential backoff for transient failures ✓
- ☐ Immediate system shutdown for error handling
- ☒ Circuit breaker patterns for external service calls ✓
- ☐ Immediate failure propagation to users with verbose logging

Show Explanation

Question: 72

CertyIQ

You're managing an agentic AI responsible for customer support ticket triage. The agent has been consistently accurate in routing tickets to the appropriate departments. However, a team leader has noticed a significant increase in the number of tickets requiring "escalation" -cases where the agent initially misclassified a complex issue as a simple, routine one, leading to delays and frustrated customers. What would be an appropriate first step in resolving this issue?

- ☒ Analyzing the agent's decision-making process, focusing on the specific criteria it uses to classify tickets, and identifying potential biases or blind spots. ✓
- ☐ Adjusting the agent's reward function to prioritize speed of resolution over accuracy, as a first step in analysis of the problem.
- ☐ Conducting a "red-teaming" exercise, having human agents deliberately create complex and ambiguous scenarios to analyze the agent's robustness.
- ☐ Increasing the agent's autonomy, granting it more decision-making power during triage to improve its efficiency.

Show Explanation

Question: 73

CertyIQ

Your team has built an agent using LangChain and needs to implement guardrails for deployment in a production environment. Which approach represents the MOST effective integration of NVIDIA NeMo Guardrails?

- ☒ Wrap the LangChain agent with NeMo Guardrails configuration while maintaining the existing workflow architecture and preserving current development investments. ✓
- ☐ Run the LangChain agent in parallel with NeMo Guardrails, allowing comparison of outputs between systems for comprehensive safety validation and performance optimization.
- ☐ Rebuild the agent using only NeMo Guardrails, thereby reconstructing the LangChain implementation with enhanced safety controls and production-ready guardrail integration.
- ☐ Configure input filtering to address safety requirements, integrating guardrail mechanisms focused on data validation and moderation within the current framework.

Show Explanation

Question: 74

CertyIQ

When designing tool integration for an agent that needs to perform mathematical calculations, web searches, and API calls, which architecture pattern provides the most scalable and maintainable approach?

- ☒ Microservice-based tool architecture with standardized interfaces ✓
- ☐ Monolithic tool handler with conditional logic for different tool types
- ☐ External tool services with manual configuration for each agent instance
- ☐ Embedded tool functions within the main agent code

Show Explanation

Question: 75

CertyIQ

When designing complex agentic workflows that include both sequential and parallel task execution, which orchestration pattern offers the greatest flexibility?

- ☐ Linear pipeline orchestration with a fixed task sequence

- ☐ Event-driven orchestration that triggers tasks reactively, in series or in parallel
- ☒ Graph-based workflow orchestration incorporating conditional branches ✓

Show Explanation

Question: 76

CertyIQ

Your team has deployed a generative agent for internal HR use, including summarizing candidate resumes and suggesting interview questions. After deployment, you've noticed that the model occasionally associates certain names or genders with particular roles. Which mitigation strategy is the most effective and scalable for reducing this type of bias in agent outputs?

- ☐ Randomly replace names in prompts to reduce identity correlation
- ☒ Implement guardrails to prevent outputs referencing protected attributes ✓
- ☐ Add more training examples to the training dataset and re-train the model
- ☐ Adjust system prompts to explicitly instruct the agent to avoid assumptions based on demographic features

Show Explanation

Question: 77

CertyIQ

This question addresses important concerns in the field of AI ethics and compliance, particularly as organizations develop more autonomous AI agents. Implementing effective guardrails against bias, ensuring data privacy, and adhering to regulations are essential components of responsible AI development. Which of the following statements accurately describes how RAGAS (Retrieval Augmented Generation Assessment) can be utilized for implementing safety checks and guardrails in agentic AI applications?

- ☐ RAGAS cannot evaluate all safety aspects independently but provides metrics like Topic Adherence and Agent Goal Accuracy that serve as guardrails. ✓
- ☒ RAGAS can only evaluate the quality of document retrieval but has no applications for safety guardrails in agentic systems. ✗
- ☐ RAGAS can only be used in conjunction with other guardrail frameworks like NeMo and cannot function independently.
- ☐ RAGAS is exclusively designed for hallucination detection and cannot evaluate other safety aspects of agentic applications.

Show Explanation

Question: 78

CertyIQ

You are rolling out a multimodal conversational agent on NVIDIA's stack: the model is containerized as a TensorRT-LLM engine, served via Triton Inference Server behind NIM microservices for routing and scaling, and protected by NeMo Guardrails for safety and compliance. During early testing, end-to-end latency exceeds your target budget, and you need to tune batching, model precision, and guardrail checks while maintaining both throughput and enforcement of safety policies. Which configuration change is most effective for reducing latency under these constraints while still enforcing NeMo Guardrails policies?

- ☐ Keep FP32 precision, increase batch size aggressively, and perform Guardrails checks in a downstream microservice after inference.
- ☐ Deploy separate Triton servers for model inference and guardrail validation, routing requests sequentially and merging outputs at the application layer.
- ☐ Quantize the TensorRT-LLM engine to INT8, disable dynamic batching, and invoke Guardrails checks synchronously within the inference path.
- ☒ Quantize the TensorRT-LLM engine to FP16, tune Triton's dynamic batching, and integrate NeMo Guardrails alongside inference to run policy checks in parallel. ✓

Show Explanation

Question: 79

CertyIQ

An AI Engineer has deployed a multi-agent system to manage supply chain logistics. Stakeholders request greater insight into how the agents decide on actions across tasks.
Which approach would best improve decision transparency without modifying the underlying model architecture?

- ☐ Generate visual summaries of attention patterns for every decision
- ☐ Gather structured user evaluations after each completed subtask
- ☒ Record a step-by-step reasoning log throughout each agent workflow ✓
- ☐ Retain and share the full sequence of task instructions with stakeholders

Show Explanation

Question: 80

CertyIQ

You are evaluating your RAG pipeline. You notice that the LLM-as-a-Judge consistently assigns high similarity scores to responses that contain irrelevant information.
What should you investigate as the most likely potential cause with the least development effort?

- ☐ The quality of the synthetic questions used for evaluation.
- ☐ The size of the knowledge base used to power the RAG pipeline.
- ☒ The prompt used to instruct the LLM-as-a-Judge to assess the response. ✓
- ☐ The temperature setting used by the LLM during response generation.

Show Explanation



Future is Secured
100% Pass Guarantee



24/7 Customer Support
Mail us - support@certyiq.com



Verified Updates
Lifetime Updates!